

Georgios Vavouliotis

✉ gvavou5@gmail.com • [gvavou5.github.io](https://github.com/gvavou5) • [in georgios-vavouliotis](https://www.linkedin.com/in/georgios-vavouliotis)

Research Interests

Computer Architecture, High Performance Computing, Machine Learning and Artificial Intelligence for Computer Architecture, Data Centers; identify and exploit the predictability of programs to co-design microarchitectural prefetching and prediction mechanisms for the memory hierarchy, re-think cache management for emerging application domains, leverage ML/AI algorithms to design intelligent microarchitectural components, and develop novel microarchitectural designs tailored to the needs of server and datacenter applications.

Education

- **Doctor of Philosophy (Ph.D.)** **Barcelona**
Universitat Politècnica de Catalunya (UPC) Sept 2018–Dec 2022
Dissertation: "[Advanced Hardware Prefetching in Virtual Memory Systems](#)" (Cum Laude Award)
Advisors: Dr. Marc Casas, Dr. Lluc Alvarez
- **Diploma on Electrical & Computer Engineering (M.Sc. equivalent)** **Athens**
National Technical University of Athens (NTUA), GPA: 8.9/10 Sept 2012–March 2018
Diploma Thesis: "[Performance Analysis of TLB Prefetching Mechanisms](#)"
Advisors: Dr. Vasileios Karakostas, Dr. Georgios Goumas

Employment

- **Senior Researcher (CPU Microarchitecture)** **Zurich**
Huawei Technologies Switzerland AG, Zurich Research Center April 2024–now
 - Performance modeling of microarchitectural features using in-house cycle-accurate simulation infrastructure to guide optimizations for next-generation Huawei cores.
 - Led the development and contributed to the integration of one (2) performance feature and one (1) power feature in Huawei CPU cores (servers, mobile).
 - Collaborating with research and product teams in the UK, Israel, and China.
 - Chartering research projects and ensuring alignment with product teams.
 - Leading the collaboration with Barcelona Supercomputing Center on synergistic cache management and cache prefetching strategies for emerging applications.
 - Authored papers published at HPCA'25, ASPLOS'25, ISCA'26, MICRO'26 and filed multiple patents on hardware prefetching and cache management as lead inventor.
 - Supervisor of interns and research engineers
 - Coordinating position openings and conducting interviews for hiring local team members and interns.
- **Postdoctoral Researcher (CPU Microarchitecture)** **Zurich**
Huawei Technologies Switzerland AG, Zurich Research Center Feb 2023–Feb 2024
 - Led in-depth research on hardware prefetching and its interaction with the virtual-memory subsystem, using both internal cycle-accurate simulators and open-source frameworks.
 - Analyzed prefetching models and coordination schemes across Huawei products, driving proposals for next-generation prefetcher designs in Huawei cores.
 - Led collaboration with Barcelona Supercomputing Center on ML-based hardware prediction and memory management for graph-processing applications.
 - Co-authored two papers published at HPCA 2024 and IPDPS 2024 and filed a patent on off-chip prediction.
- **Research Intern (Memory Management)** **Zurich**
Huawei Technologies Switzerland AG, Zurich Research Center Nov 2021–April 2022
 - Researched the impact of Concurrent Java Garbage Collectors (e.g., ZGC) on cache locality of modern microarchitectures
 - Authored a paper published at ISMM 2023 [[Best Paper Award 🏆](#)]
- **Researcher (Computer Architecture & Operating Systems)** **Barcelona**
Barcelona Supercomputing Center (BSC-CNS) Sept 2018–Dec 2022
 - Carried out research on address translation, speculation mechanisms for microarchitecture, and interaction between computer architecture and artificial intelligence

- Contributed to the writing of research proposals for international grants
- Contributed to collaborative projects on cache management and off-chip prediction

Honors and Awards

- **HiPEAC 2025 Paper Award & Skyline Award (by Huawei)** 🏆 Zurich
For the [To Cross, or Not to Cross Pages for Prefetching?](#) paper 2025
- **HiPEAC 2025 Paper Award & Skyline Award (by Huawei)** 🏆 Zurich
For the [Instruction-Aware Cooperative TLB and Cache Replacement Policies](#) paper 2025
- **Innovation Award** 🏆 Zurich
For contributing to the development of a performance feature landed in Huawei cores 2025
- **Future Star Award** 🏆 Zurich
For my research contributions to microarchitectural designs of Huawei CPU cores 2025
- **HiPEAC 2024 Paper Award & Skyline Award (by Huawei)** 🏆 Zurich
For the [A Two Level Neural Approach Combining Off-Chip Prediction with Adaptive Prefetch Filtering](#) paper 2024
- **Core Star Award (by Huawei)** 🏆 Zurich
For the [Concurrent GCs and Modern Java Workloads: A Cache Perspective](#) paper 2023
- **Cum Laude Ph.D. Award** 🏆 Spain
For my Ph.D. dissertation entitled *Advanced Hardware Prefetching in Virtual Memory Systems* 2023
- **Best Paper Award - ISMM 2023** 🏆 Orlando
For the [Concurrent GCs and Modern Java Workloads: A Cache Perspective](#) paper 2023
- **HiPEAC 2022 Paper Award** 🏆 Chicago
For the [Page Size Aware Cache Prefetching](#) paper 2022
- **HiPEAC 2021 Paper Award** 🏆 Virtual
For the [Morrigan: A Composite Instruction TLB Prefetcher](#) paper 2021
- **HiPEAC 2021 Paper Award** 🏆 Virtual
For the [Exploiting Page Table Locality for Agile TLB Prefetching](#) paper 2021
- **Best Poster Award** 🏆 Virtual
2021 ACM School on HPC Computer Architectures for AI and Dedicated Applications 2021
- **2019 FPI Doctoral Fellowship** 🏆 Barcelona
Funded by Spanish Government (MINECO) 2019-2022
- **Mathematics Award - National Technical University of Athens** 🏆 Athens
Excellent marks in all mathematical courses during the first two years. 2014
- **"The Great Moment of Education" Eurobank EFG Scholarship** 🏆 Chalkida
Highest rank in National Qualifications Exams in my school 2012

Peer-Reviewed Conference Publications

- [1] **Georgios Vavouliotis**, Damjan Lampret, Sara Friedman, Leor Peled, [Exploiting Program Semantics for Hardware Prefetching](#). Proceedings of the 59th IEEE/ACM International Symposium on Microarchitecture (MICRO), Athens, Nov 2022.
- [2] Diamantis Patsidis, Leor Peled, **Georgios Vavouliotis**, [Enabling Adaptive Policy Arbitration with Context Awareness](#). Proceedings of the 59th IEEE/ACM International Symposium on Microarchitecture (MICRO), Athens, Nov 2022.
- [3] **Georgios Vavouliotis**, Tom Rollet, Davide Basilio Bartolini, Boris Grot, Leor Peled, Yang Lixia, [Bumper: Hinting Instruction Usefulness for Robust Unified Caches](#). Proceedings of the International Symposium on Computer Architecture (ISCA), Raleigh, June 2026.
- [4] Alexandre Valentin Jamet, **Georgios Vavouliotis**, Martí Torrents, Dimitrios Chasapis, Marc Casas, [Enhancing Instruction Prefetching via Cache and TLB Management](#). Proceedings of the International Symposium on Computer Architecture (ISCA), Raleigh, June 2026.
- [5] **Georgios Vavouliotis**^{*}, Dimitrios Chasapis^{*}, Daniel A. Jiménez, Marc Casas, [Instruction-Aware Cooperative TLB and Cache Replacement Policies](#). Proceedings of the International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), Rotterdam, April 2025.
- [6] **Georgios Vavouliotis**, Marti Torrents, Boris Grot, Kleovoulos Kalaitzidis, Leor Peled, Marc Casas, [To Cross,](#)

or Not to Cross Pages for Prefetching?. Proceedings of the 31st International Symposium on High-Performance Computer Architecture (HPCA), Las Vegas, March 2025.

[7] Alexandre Valentin Jamet, **Georgios Vavouliotis**, Lluc Alvarez, Daniel A. Jiménez, Marc Casas, [Practically Tackling Memory Bottlenecks of Graph-Processing Workloads](#). Proceedings of the 38th IEEE International Parallel & Distributed Processing Symposium (IPDPS), San Francisco, May 2024.

[8] Alexandre Valentin Jamet, **Georgios Vavouliotis**, Lluc Alvarez, Daniel A. Jiménez, Marc Casas, [A Two Level Neural Approach Combining Off-Chip Prediction with Adaptive Prefetch Filtering](#). Proceedings of the 30th International Symposium on High-Performance Computer Architecture (HPCA), Edinburgh, March 2024.

[9] Maria Carpen-Amarie, **Georgios Vavouliotis**, Konstantinos Tovletoglou, Boris Grot, Rene Mueller, [Concurrent GCs and Modern Java Workloads: A Cache Perspective](#). Proceedings of the 2023 ACM SIGPLAN International Symposium on Memory Management (ISMM), Orlando, June 2023. [Best Paper Award 🏆]

[10] **Georgios Vavouliotis**, Gino Chancon, Lluc Alvarez, Paul V. Gratz, Daniel A. Jiménez, Marc Casas, [Page Size Aware Cache Prefetching](#). Proceedings of the 55th IEEE/ACM International Symposium on Microarchitecture (MICRO), Chicago, Oct 2022.

[11] **Georgios Vavouliotis**, Lluc Alvarez, Boris Grot, Daniel A. Jiménez, Marc Casas, [Morrigan: A Composite Instruction TLB Prefetcher](#). Proceedings of the 54th IEEE/ACM International Symposium on Microarchitecture (MICRO), Virtual, Oct 2021.

[12] **Georgios Vavouliotis**, Lluc Alvarez, Vasileios Karakostas, Konstantinos Nikas, Nectarios Koziris, Daniel A. Jiménez, Marc Casas, [Exploiting Page Table Locality for Agile TLB Prefetching](#). Proceedings of the 48th International Symposium on Computer Architecture (ISCA), Virtual, June 2021. Also presented as a [Poster](#) at the 49th International Symposium on Computer Architecture (ISCA), New York, June 2022.

Peer-Reviewed Journal Publications

[1] Diamantis Patsidis, **Georgios Vavouliotis**, [Context-Aware Set Dueling for Dynamic Policy Arbitration](#), IEEE Computer Architecture Letters (CAL), Oct 2025.

Peer-Reviewed Workshop Publications and Posters

[1] **Georgios Vavouliotis**, Gino Chancon, Lluc Alvarez, Paul V. Gratz, Daniel A. Jiménez, Marc Casas, [Leveraging Page Size Information to Enhance Data Cache Prefetching](#). Poster at the 2021 ACM Summer School on HPC Computer Architectures for AI and Dedicated Applications, Virtual, Aug 2021. [Best Poster Award 🏆]

[2] **Georgios Vavouliotis**, Lluc Alvarez, Marc Casas, [Pushing the envelope on free TLB prefetching](#). Proceedings of the 8th BSC Doctoral Symposium, Virtual, May 2021.

[3] **Georgios Vavouliotis**, [Cost-Effective Instruction TLB Prefetching](#). Second Young Architect Workshop (YArch 2020). In conjunction with The 25th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), Virtual, March 2020.

Patents

[1] "PATENT ON HARDWARE PAGE TABLE WALKING" [pending]

[2] "PATENT ON CACHE MANAGEMENT" [pending]

[3] "Method and corresponding apparatus for filtering prefetch requests", Georgios Vavouliotis, Boris Grot, Leor Peled, Filed: 20/06/2024

[4] "PERCEPTRON-BASED OFF-CHIP PREDICTOR", Alexandre Valentin Jamet, **Georgios Vavouliotis**, Marc Casas, Filed: 21/12/2023

Dissertations

[1] **Georgios Vavouliotis** "[Advanced Hardware Prefetching in Virtual Memory Systems](#)". Book of my Doctoral dissertation. [Cum Laude Award 🏆]

Thesis Statement: *Hardware TLB prefetching can reduce the address translation overheads posed with applications with large data and code footprints while exploiting address translation metadata available at the microarchitecture and runtime levels can improve the performance of cache prefetchers.*

[2] **Georgios Vavouliotis**, Goumas Georgios, "[Performance Analysis of TLB Prefetching Mechanisms](#)". Book of my Diploma thesis (in Greek).

Invited Talks

- **Interaction between Virtual Memory and Hardware Prefetching** **Seoul**
Virtuoso Workshop (MICRO 2025) *October 2025*
- **To Cross, or Not to Cross Pages for Prefetching?** **Las Vegas**
2025 IEEE International Symposium on High-Performance Computer Architecture (HPCA) *Feb 2025*
- **Importance of Single-Thread Performance in the Many-Core Era** **Buenos Aires**
SPARKS Workshop (ISCA 2024) *July 2024*
- **Microarchitectural Prediction and Prefetching in Virtual Memory Systems** **Online**
AMD Research *May 2024*
- **Advanced Hardware Prefetching in Virtual Memory Systems** **Athens**
6th Computing Systems Research Day @ CSLab *Jan 2023*
- **Page Size Aware Cache Prefetching** **Chicago**
55th IEEE/ACM International Symposium on Microarchitecture (MICRO) *Oct 2022*
- **Morrigan: A Composite Instruction TLB Prefetcher** **Zurich**
Huawei Zurich Research Center *Dec 2021*
- **Morrigan: A Composite Instruction TLB Prefetcher** **Virtual**
54th IEEE/ACM International Symposium on Microarchitecture (MICRO) *Oct 2021*
- **Leveraging Page Size Information to Enhance Data Cache Prefetching** **Virtual**
ACM School on HPC Computer Architectures for AI and Dedicated Applications *Aug 2021*
- **Exploiting Page Table Locality for Agile TLB Prefetching** **Virtual**
48th International Symposium on Computer Architecture (ISCA) *June 2021*
- **Pushing the Envelope on Free TLB Prefetching** **Virtual**
8th BSC Doctoral Symposium *May 2021*

Technical Reviewer for Conferences and Journals

- International Symposium on High-Performance Computer Architecture (HPCA'27)
- International Symposium on Computer Architecture (ISCA'26)
- International Symposium on Advanced Parallel Processing Technology (APPT'26)
- IEEE Computer Architecture Letters (CAL'26)
- IEEE Transactions on Computers (TC'26)
- ACM Transactions on Architecture and Code Optimization (TACO'26)
- ACM Transactions on Modeling and Performance Evaluation of Computing Systems (TOMPECS'26)
- International Symposium on High-Performance Computer Architecture (HPCA'26)
- IEEE Transactions on Network Science and Engineering (TNSM'25)
- IEEE Transactions on Software Engineering (ToS'25)
- IEEE Micro
- ACM Transactions on Architecture and Code Optimization (TACO'25)
- International Symposium on Microarchitecture (MICRO'25)
- International Symposium on Computer Architecture (ISCA'25)
- International Symposium on High-Performance Computer Architecture (HPCA'25)
- IEEE Micro
- International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS'25)
- International Symposium on Microarchitecture (MICRO'24)
- International Symposium on Computer Architecture (ISCA'24)
- ACM Transactions on Architecture and Code Optimization (TACO'24)
- European Conference on Parallel and Distributed Computing (Euro-Par'24)
- International Symposium on High-Performance Computer Architecture (HPCA'24, Artifact Evaluation)
- Journal of Parallel and Distributed Computing (JPDC'23)
- IEEE Access 2023
- ACM Transactions on Architecture and Code Optimization (TACO'23)

- IEEE Supercomputing Conference (SC'21)
- IEEE Transactions on Computers (2021)
- ACM International Conference on Computing Frontiers (CF'21)

Professional Activities - Leadership & Service

- ISCA'24 **Buenos Aires**
2024
 - *Website Chair*

Teaching Experience

- **Private Tutor** **Athens**
2013–2018
 - *C programming, Computer Architecture, Automatic Control for undergraduate students*
- **Lab Assistant at National Technical University of Athens (NTUA)** **Athens**
2013-2017
 - *C programming*

Selected Press

- **TCuArch Interview** **youtube**
Dec 2022
 - *What attracted you to research computer architecture and systems?*
- **Page Size Aware Cache Prefetching** **easypref.net**
Oct 2022
 - *Oct Newsletter of easypref.net – Among the Most Interesting MICRO '22 Papers*

Languages

- Greek (Native)
- English
- Italian (Beginner)
- Spanish (Beginner)